

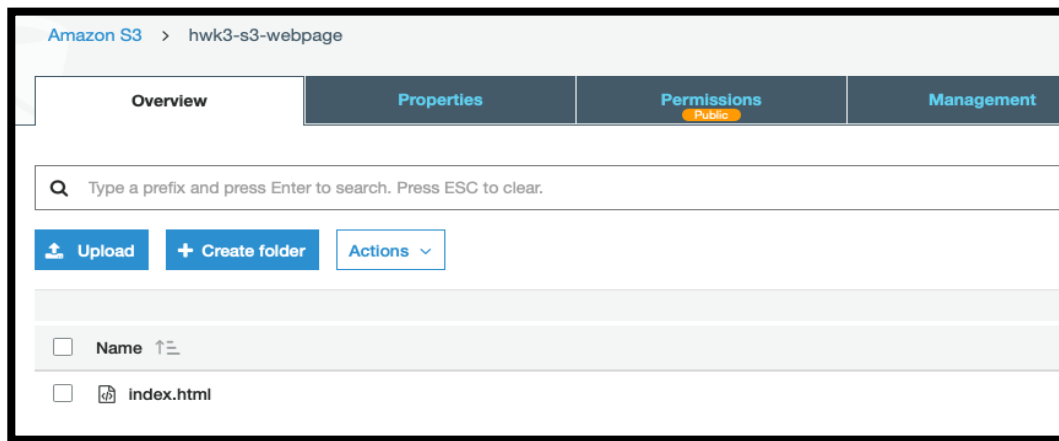
E91 Cloud Devops: Fall 2018
Assignment 3
Stephen Akaeze
https://code.harvard.edu/sta283/cscie-91_sta283

Assignment Instruction manual:

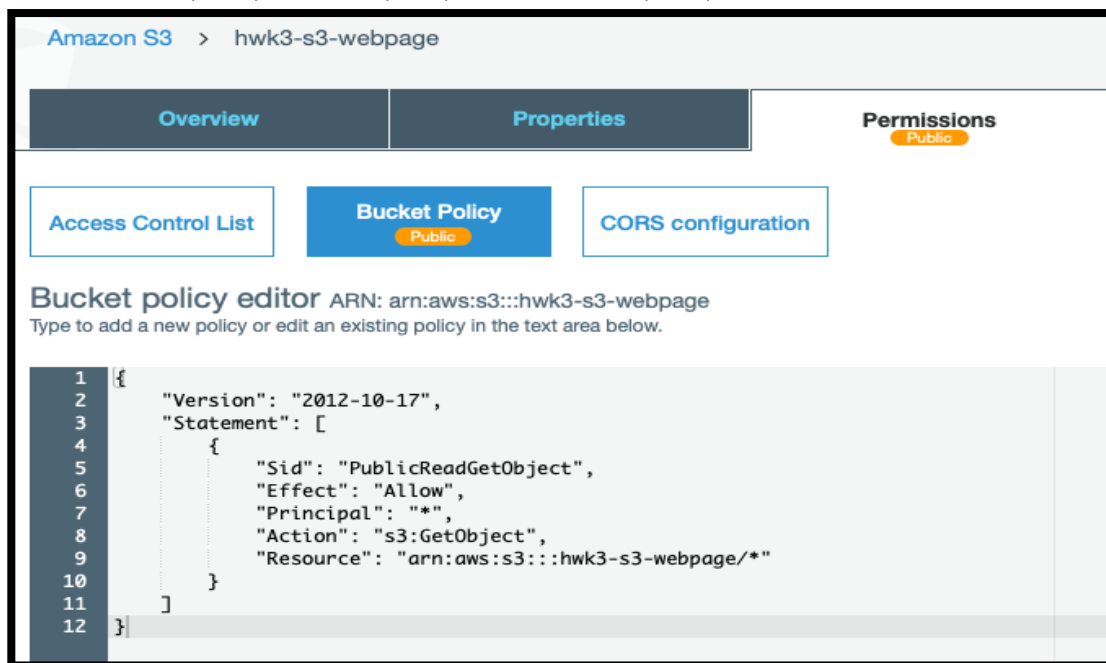
<https://canvas.harvard.edu/courses/53026/files/6670661/download?verifier=ULySloPaZMkA0SRkueEVoHrmTzHWwfBhzVslrj2R&wrap=1>

Problem 1:

- 1) The S3 bucket was created and named "hwk3-s3-webpage". The S3 bucket's property was set to host a static website. An index.html file containing two lines was created and uploaded to webpage.



- 2) The S3 bucket policy was setup as public with the policy shown below



- 3) The site was checked and it works. The site URL is <http://hwk3-s3-webpage.s3-website-us-west-1.amazonaws.com>

Problem 2

- 1) An instance with AWS Linux AMI was bootstrapped with the user data below and created in the public subnet created in homework 2. The ssh keys and new user implementation are also included in user data. Apache HTTP server was installed with the line "Assignment 3: Public Web Server" as also show below

```
apache_bootstrap.sh
1  #!/bin/bash
2
3  yum install httpd -y
4  echo "Assignment 3: Public Web Server" > /var/www/html/index.html
5
6  systemctl start httpd
7  systemctl enable httpd
8
9  userdel -r stakes 2> /dev/null
10 useradd stakes --create-home --shell /bin/bash
11
12 mkdir /home/stakes/.ssh
13
14 cat > /home/stakes/.ssh/authorized_keys << EOF
15 ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQDdyS33YJTTooUWC6VdcM7A0SPAGRCuIaXFsp7EL336NYL
16 EOF
17
18 chown -R stakes.stakes /home/stakes/.ssh
19 chmod 700 /home/stakes/.ssh
20 chmod 600 /home/stakes/.ssh/authorized_keys
21
22 echo "stakes ALL=(ALL) NOPASSWD:ALL" > /etc/sudoers.d/stakes
```

The EC2 instance(hwk3_vpc1_pub) was created and running

Instance State: Running

Add filter

Name

Instance ID

Instance Type

Availability Zone

Instance State

Status Checks

Alarm Status

Public IP

IPv4 Public IP

IPv6 IPs

Key Name

hwk3_vpc1_pub

i-0654c75ce63e10712

t2.micro

us-west-1a

running

2/2 checks ...

None

54.193.73.237

-

Description

Status Checks

Monitoring

Tags

Instance ID

i-0654c75ce63e10712

Public DNS (IPv4)

-

Instance state

running

IPv4 Public IP

54.193.73.237

Instance type

t2.micro

IPv6 IPs

-

Elastic IPs

Private DNS

ip-192-168-1-180.us-west-1.compute.internal

Availability zone

us-west-1a

Private IPs

192.168.1.180

Security groups

[hwk2_ssh_internet](#), [hwk2_internet_internet](#). [view inbound rules](#). [view outbound rules](#)

Secondary private IPs

Scheduled events

[No scheduled events](#)

VPC ID

vpc-00307c7ffde8a31ac

AMI ID

[amzn2-ami-hvm-2.0.20181024-x86_64-gp2 \(ami-01beb64058d271bc4\)](#)

Subnet ID

subnet-0335671bcd1625460

Platform

-

Network interfaces

[eth0](#)

IAM role

-

Source/dest. check

True

Key pair name

-

T2/T3 Unlimited

Disabled

Owner

293836347668

EBS-optimized

False

Launch time

November 3, 2018 at 12:06:19 AM UTC-7 (less than one hour)

Root device type

ebs

Termination protection

False

Root device

[/dev/xvda](#)

Lifecycle

normal

Block devices

[/dev/xvda](#)

The webpage was checked and the apache server was running

```
stakes@CSCI91:~$ curl 54.193.73.237:80
Assignment 3: Public Web Server
```

Note: the assignment 2 private subnet was replaced with a new subnet in AZ **us-west-1b** While the public subnet remained in AZ **us-west-1a**. Thus, having two availability zones in the original VPC

The same above steps was used to create another instance in the private subnet in different AZ. The private instance apache server is only serving “Assignment 3: private Web Server”. The instance was launched and it was working as shown below,

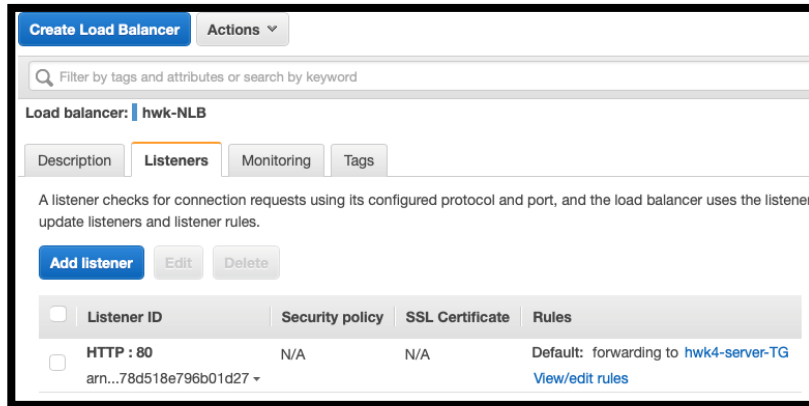
Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public IP	IPv4 Public IP	IPv6 IPs	Key Name
hwk3_vpc1_priv	i-027331ceebc1553cf	t2.micro	us-west-1b	running	2/2 checks ...	None	-	-	-	-

Instance: i-027331ceebc1553cf (hwk3_vpc1_priv) Private IP: 192.168.2.166	
Description	Status Checks Monitoring Tags
Instance ID	i-027331ceebc1553cf
Instance state	running
Instance type	t2.micro
Elastic IPs	
Availability zone	us-west-1b
Security groups	hwk2_ssh_internet, hwk2_ssh_publicSubnet, hwk2_internet_internet. view inbound rules. view outbound rules
Scheduled events	No scheduled events
AMI ID	amazon2-ami-hvm-2.0.20181024-x86_64-gp2 (ami-01beb64058d271bc4)
Platform	-
IAM role	-
Key pair name	-
Owner	293836347668
Launch time	November 3, 2018 at 12:20:30 AM UTC-7 (less than one hour)
Termination protection	False
Lifecycle	normal
Public DNS (IPv4)	-
IPv4 Public IP	-
IPv6 IPs	-
Private DNS	ip-192-168-2-166.us-west-1.compute.internal
Private IPs	192.168.2.166
Secondary private IPs	
VPC ID	vpc-00307c7ffde8a31ac
Subnet ID	subnet-0473a1a37e715641d
Network interfaces	eth0
Source/dest. check	True
T2/T3 Unlimited	Disabled
EBS-optimized	False
Root device type	ebs
Root device	/dev/xvda
Block devices	/dev/xvda

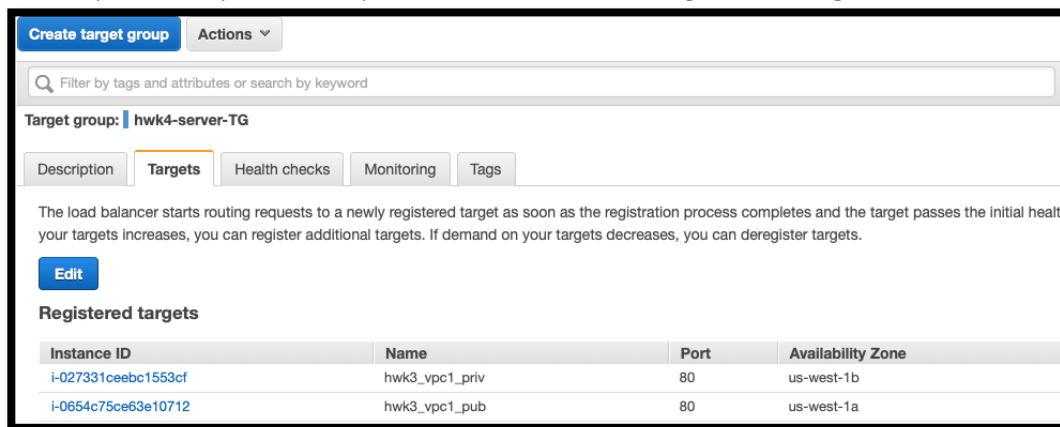
Using ssh access to the public instance, the private server was verified as working

```
[stakes@ip-192-168-1-180 ~]$ curl 192.168.2.166:80
Assignment 3: Private Web Server
```

7) The network load balancer created using the application load balancer. It was placed in the HW2 VPC with Web port 80 open to incoming traffic. **DNS Name is hwk-NLB-718786712.us-west-1.elb.amazonaws.com**



The target group was also created as described in the instruction manual including the recently created public and private instances as its registered targets,



Loading details

To confirm that the load balancer/target group setup works, a custom load balancing script was written. The load balancer IP address was recovered using *"host <dns name>"*

With the IP address, my VM sends multi-process looped http requests to the load balancer's resolved IP via the "curl" command. **I could not get JMeter to work so I improvised.** The looped commands per process is shown below
`curl -s http://54.193.86.88:80?\[1-1000\]`

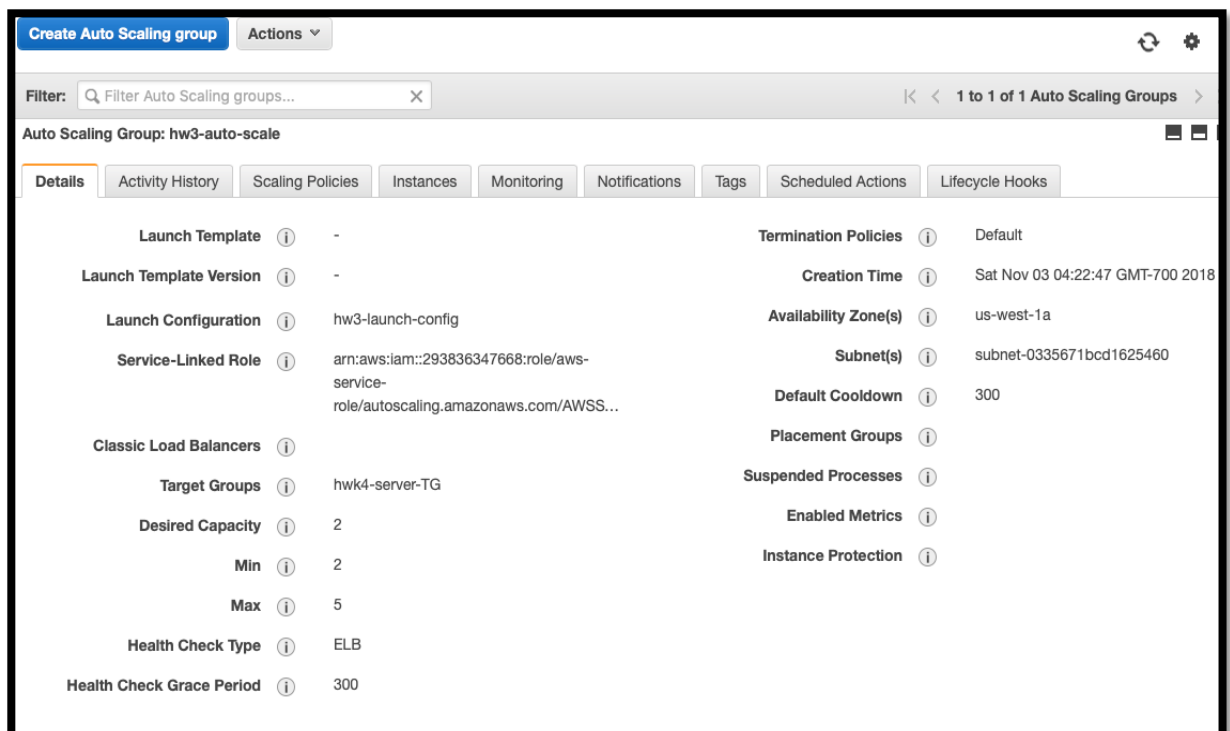
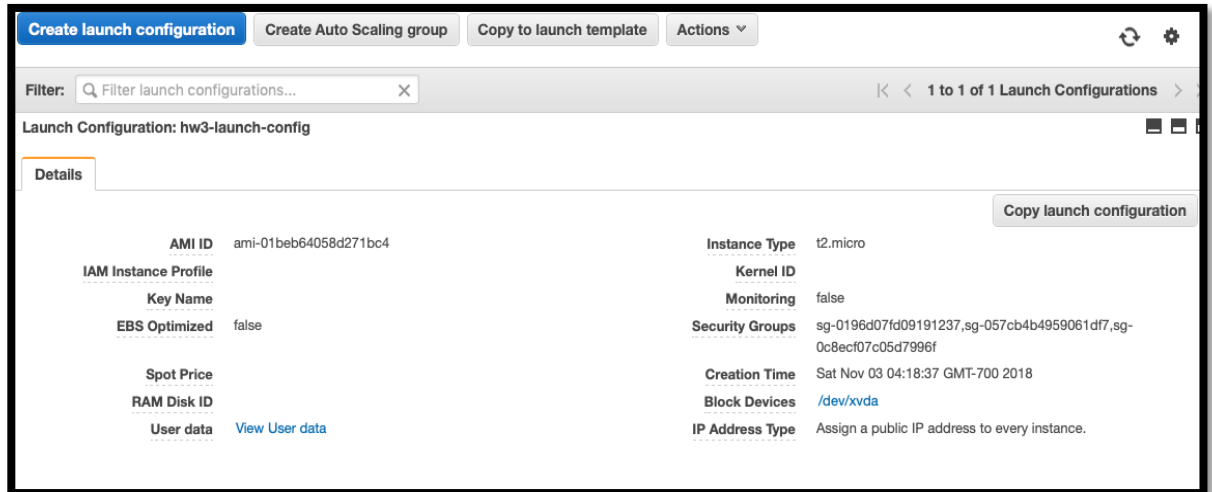
The result shows both machines responding to the load balancer's http requests

```
Damien:~ Damien$ curl -s "http://54.193.86.88:80?[1-1000]"
Assignment 3: Public Web Server
Assignment 3: Private Web Server
Assignment 3: Public Web Server
Assignment 3: Private Web Server
Assignment 3: Public Web Server
Assignment 3: Private Web Server
Assignment 3: Public Web Server
Assignment 3: Private Web Server
```

When the load goes up, the DNS name resolves to more IP addresses.

Problem 3

- 1) As specified by the instruction manual, the launch configuration and Auto Scaling were created as directed as shown below,



- a. As instructed by the instruction manual, auto scaling starts with 2 instances. When one instance was terminated, the site did not go down because the second instance was still running. The site only goes down when both instances are terminated. After the single instance termination, the instance was immediately deleted from the target group as shown below.

Two auto scale instances are running fine

Launch Instance ▼ Connect Actions ▼							
Name : Auto Scale Instance ✕ Add filter							
<input type="checkbox"/>	Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status
<input type="checkbox"/>	Auto Scale Instance	i-04e2950fadcce86e5	t2.micro	us-west-1a	terminated		None
<input type="checkbox"/>	Auto Scale Instance	i-054a3d4bdb7494f99	t2.micro	us-west-1a	terminated		None
<input type="checkbox"/>	Auto Scale Instance	i-0880e2cd7ffb92632	t2.micro	us-west-1a	running	2/2 checks ...	None
<input type="checkbox"/>	Auto Scale Instance	i-0b4e99323ab2ce6e4	t2.micro	us-west-1a	running	2/2 checks ...	None

Target group shows both running instances as healthy

Registered targets				
Instance ID	Name	Port	Availability Zone	Status
i-0880e2cd7ffb92632	Auto Scale Instance	80	us-west-1a	healthy
i-027331ceebc1553cf	hwk3_vpc1_priv	80	us-west-1b	unused
i-0654c75ce63e10712	hwk3_vpc1_pub	80	us-west-1a	unused
i-0b4e99323ab2ce6e4	Auto Scale Instance	80	us-west-1a	healthy

Then, one instance was terminated

Launch Instance ▼ Connect Actions ▼							
Name : Auto Scale Instance ✕ Add filter							
<input type="checkbox"/>	Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status
<input type="checkbox"/>	Auto Scale Instance	i-04e2950fadcce86e5	t2.micro	us-west-1a	terminated		None
<input type="checkbox"/>	Auto Scale Instance	i-054a3d4bdb7494f99	t2.micro	us-west-1a	terminated		None
	Auto Scale Instance	i-0880e2cd7ffb92632	t2.micro	us-west-1a	shutting-do...		None
<input type="checkbox"/>	Auto Scale Instance	i-0b4e99323ab2ce6e4	t2.micro	us-west-1a	running	2/2 checks ...	None

The target group instantly deleted the terminated instance

Registered targets				
Instance ID	Name	Port	Availability Zone	Status
i-027331ceebc1553cf	hwk3_vpc1_priv	80	us-west-1b	unused
i-0654c75ce63e10712	hwk3_vpc1_pub	80	us-west-1a	unused
i-0b4e99323ab2ce6e4	Auto Scale Instance	80	us-west-1a	healthy

- b. The autoscaling immediately created another instance in its place and the website is now accessible from both instances

Launch Instance ▼ Connect Actions ▼							
Name : Auto Scale Instance ✕ Add filter							
<input type="checkbox"/>	Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status
<input type="checkbox"/>	Auto Scale Instance	i-04e2950fadcce86e5	t2.micro	us-west-1a	terminated		None
<input type="checkbox"/>	Auto Scale Instance	i-04f8e609a6d8a31ce	t2.micro	us-west-1a	running	Initializing	None
<input type="checkbox"/>	Auto Scale Instance	i-054a3d4bdb7494f99	t2.micro	us-west-1a	terminated		None
<input type="checkbox"/>	Auto Scale Instance	i-0880e2cd7ffb92632	t2.micro	us-west-1a	terminated		None
<input type="checkbox"/>	Auto Scale Instance	i-0b4e99323ab2ce6e4	t2.micro	us-west-1a	running	2/2 checks ...	None

- d. The Load was applied and an additional instance was created indicating that the 35% CPU load had been met by both load balancer target group instances

Problem 4:

- 1) The “dev” branch assignment 2 folder merging to “master” branch is shown below

```
stakes@CSCI91:~/cscie-91_sta283$ git checkout master
Switched to branch 'master'
Your branch is up to date with 'origin/master'.
```

```
stakes@CSCI91:~/cscie-91_sta283$ git checkout dev -- assignment-02
stakes@CSCI91:~/cscie-91_sta283$ git status
On branch master
Your branch is up to date with 'origin/master'.
Changes to be committed:
  (use "git reset HEAD <file>..." to unstage)

    new file:   assignment-02/README.md
```

```
stakes@CSCI91:~/cscie-91_sta283$ git commit -m "committing assignment-02 folder from dev to master branch"
[master 2ba41e5] committing assignment-02 folder from dev to master branch
 1 file changed, 0 insertions(+), 0 deletions(-)
 create mode 100644 assignment-02/README.md
stakes@CSCI91:~/cscie-91_sta283$ git push
Enumerating objects: 5, done.
Counting objects: 100% (5/5), done.
Compressing objects: 100% (2/2), done.
Writing objects: 100% (4/4), 389 bytes | 389.00 KiB/s, done.
Total 4 (delta 0), reused 0 (delta 0)
To code.harvard.edu:sta283/cscie-91_sta283.git
 6db50ff..2ba41e5  master -> master
```

- 2) The assignment 3 folder was created in “dev” and the documentation text file (assig3_quest3_doc.txt) was added to the assignment 3 folder and pushed to code.harvard.edu (link: https://code.harvard.edu/sta283/cscie-91_sta283/blob/dev/assignment-03/assig3_quest3_doc.txt). The documentation file also contains as URL that helps demonstrate the relationship between autoscaling and load balancers. A screenshot of the text pushed assignment 3 folder and text file is shown below,

sta283 / cscie-91_sta283 Private

Watch 0

Star 0

Fork 0

<> Code

Issues 0

Pull requests 0

Projects 0

Wiki

Insights

Settings

Branch: dev

cscie-91_sta283 / assignment-03 / assig3_quest3_doc.txt

Find file

Copy path

sta283 assignment 3 documentation file

1b96cd2 2 hours ago

1 contributor

24 lines (11 sloc) | 1.84 KB

Raw

Blame

History

1 Assignment 3 Question 3 Documentation

2

3 1) describe the monthly cost of using S3 static page vs using a single EC2 instance

4 Answer:

5

6 S3 monthly cost: The S3 static page monthly cost is cheapest because AWS charges just \$0.0044 per Get and all other requests. Y

7

8 EC2 monthly cost: The EC2 static page monthly cost is more expensive than the SP3 cost because it is charged per hour. At \$0.04

9

10 Load Balancing and Autoscaling monthly cost: This is the most expensive of the three options as it adds the \$0.0252 per applica

11

12 2) EC2s are used for websites because they can easily support dynamic sites, support autoscaling &load balancers, provide high

13

14 3) Relationship between auto scaling and load balancing is best captured by the link below

15

16 Link:https://d2908q01vomqb2.cloudfront.net/7719a1c782a1ba91c031a682a0a2f8658209adbf/2017/04/18/CDBG-topology.png

17

18 As shown in the image above

19

20 The Load Balancer Sends traffic to backsend servers(EC2 instances in a particulat target group) and also monitors their individ

21

22

23