

HARVARD
Extension School

Week 7

AWS Services - DevOps Tools

This week..

- **Week 07:** AWS Services - DevOps Tools

This lecture will discuss some of the DevOps resources, tools, and methods introduced by AWS like CloudFormation, Auto Scaling, CloudWatch, CloudTrail and Cloud Config. The focus will be on the use of Auto-scaling for high availability and scalability, and on CloudFormation to build AWS resources using CloudFormation templates.

- High availability and scalability
 - Auto Scaling
 - Elastic Load Balancer
- CloudFormation
 - Templates
 - Template Designer
 - Templates output and exports
- Monitoring and Logging
 - CloudWatch
 - CloudTrail
 - CloudConfig
 - Elasticsearch

DevOps & Agile - reminder

- DevOps values:
 - Agility: Do things faster
 - Cost Efficiency: Do things cheaper
 - Quality: Do things that will meet the needs of the business and end users.

Keys to consider

When choosing devops tools:

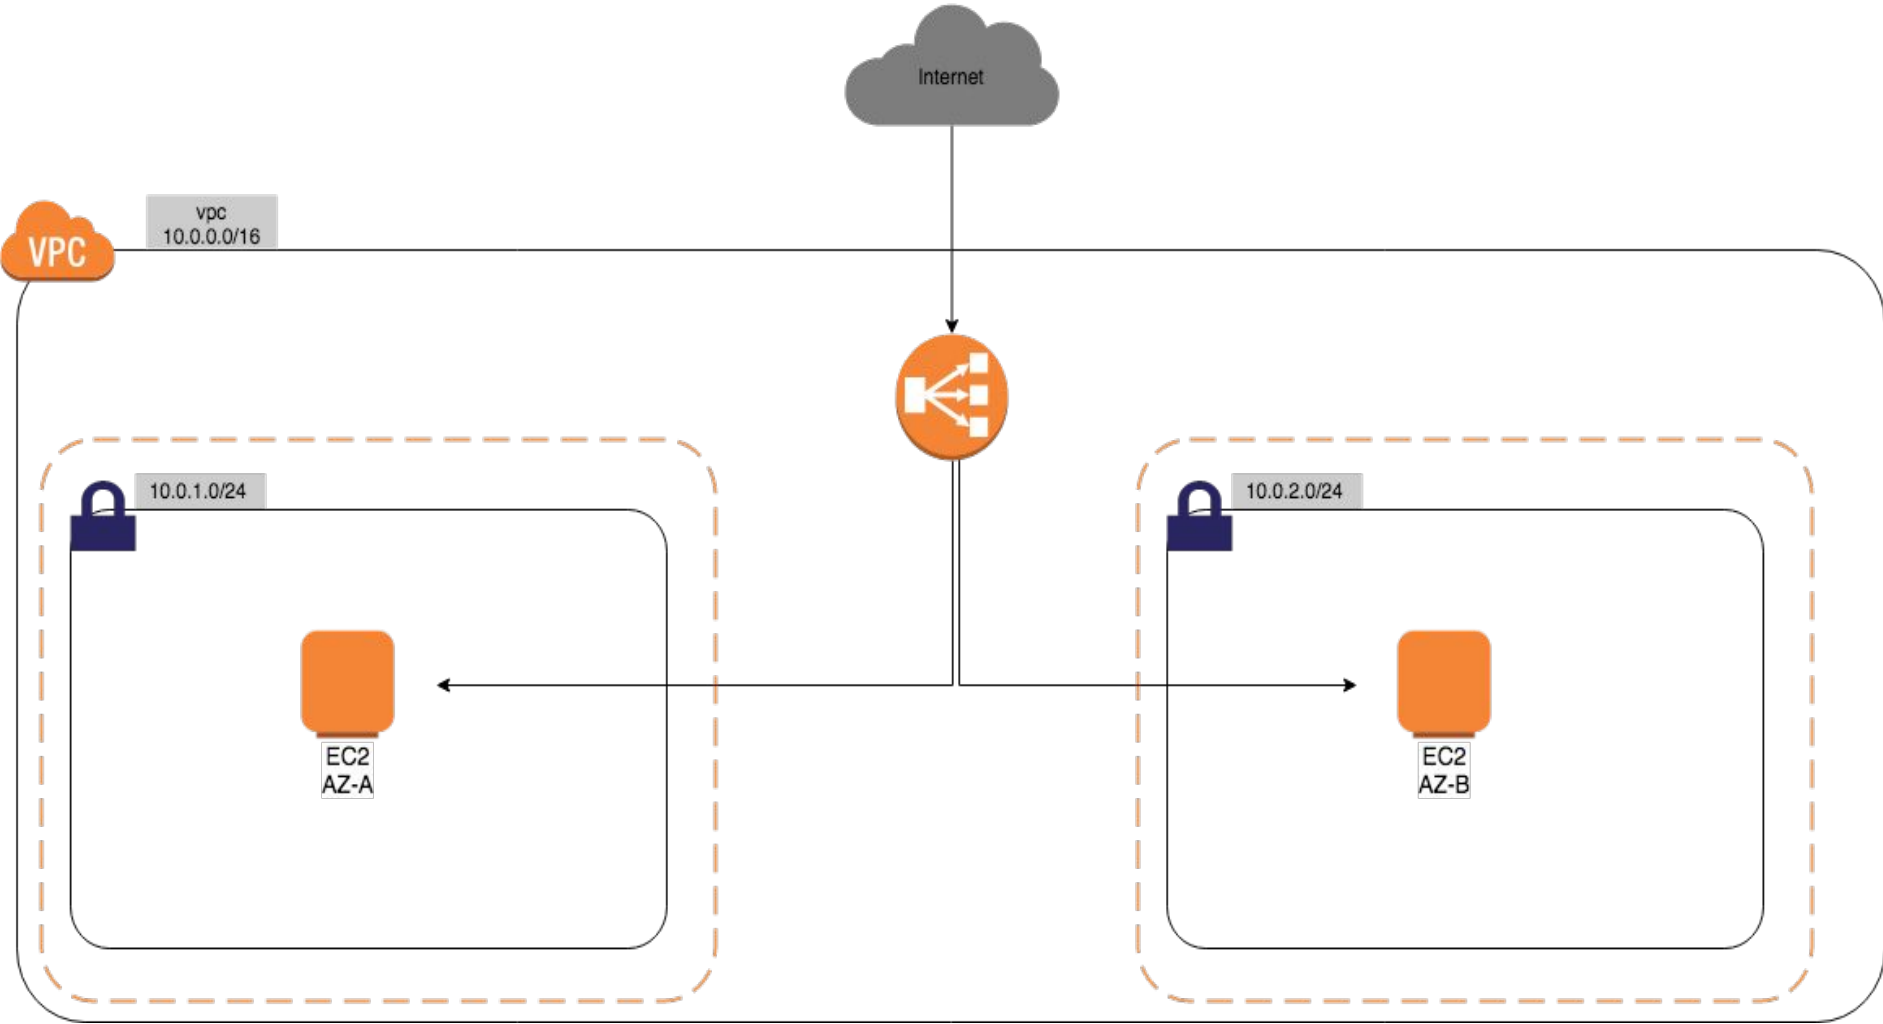
- Automate as much as possible
- Cost
- Integration with the cloud platform in use
- Developers acceptance
- Tools will be replaceable
- Not limited to services/tools provided by the cloud platform in use
- It's ok to pick the wrong tool but don't continue using the wrong tool

Testing tools process

1. Test the chosen tool as a unit
2. Test the chosen tool with other tools in place
3. Test the complete process

EC2 Load Balancers for availability

- Service based on software to functionally route a changing number of EC2 instances in different AZs within the same region.
- The Load Balancer is a single point of contact to the outside.
- Two types of Load Balancers:
 - Monitor health check at the application level (ALB - Application Load Balancers): preferred for http or https
 - Monitor at the network level (ELB - Elastic Load Balancers): preferred for http or https as well as TCP or SSL
- If any of the EC2s in a specific Target Group in any of the zones become unavailable, the Load Balancer simply re-routes traffic to the available servers.



Load Balancer

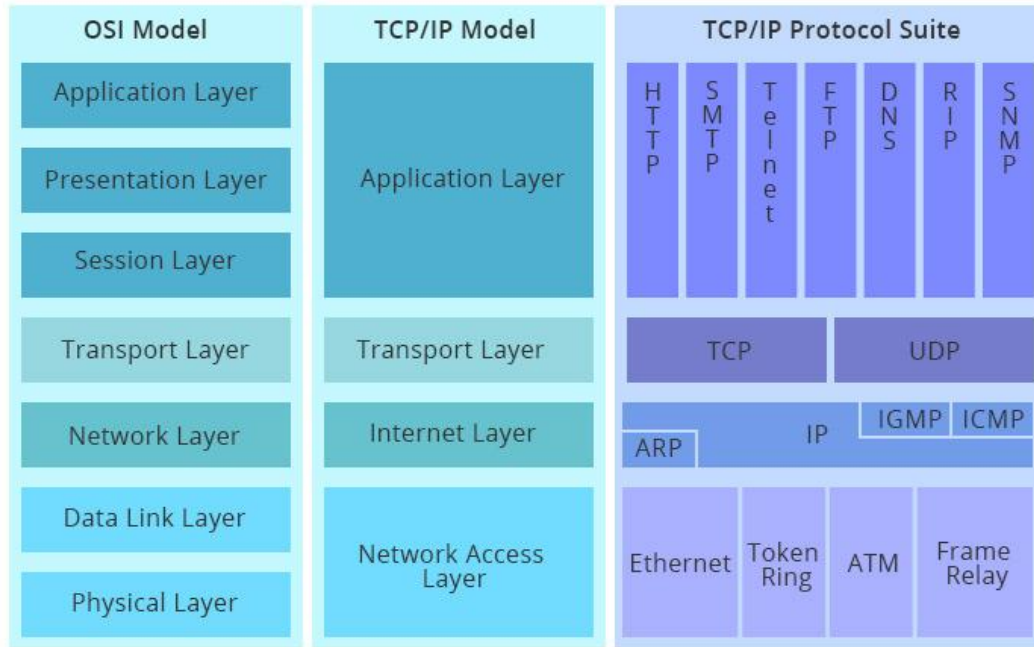
- Layer 7 HTTP/HTTPS Application Layer Load Balancer
 - ◆ Target groups and Rules to route paths

- Layer 4 Network Load Balancer
 - ◆ Extreme performance is required.
 - ◆ Low latency while handling millions of requests per second

- Classic Load Balancer: Layer 4 and can use layer 7
 - ◆ X-Forwarding
 - ◆ sticky session

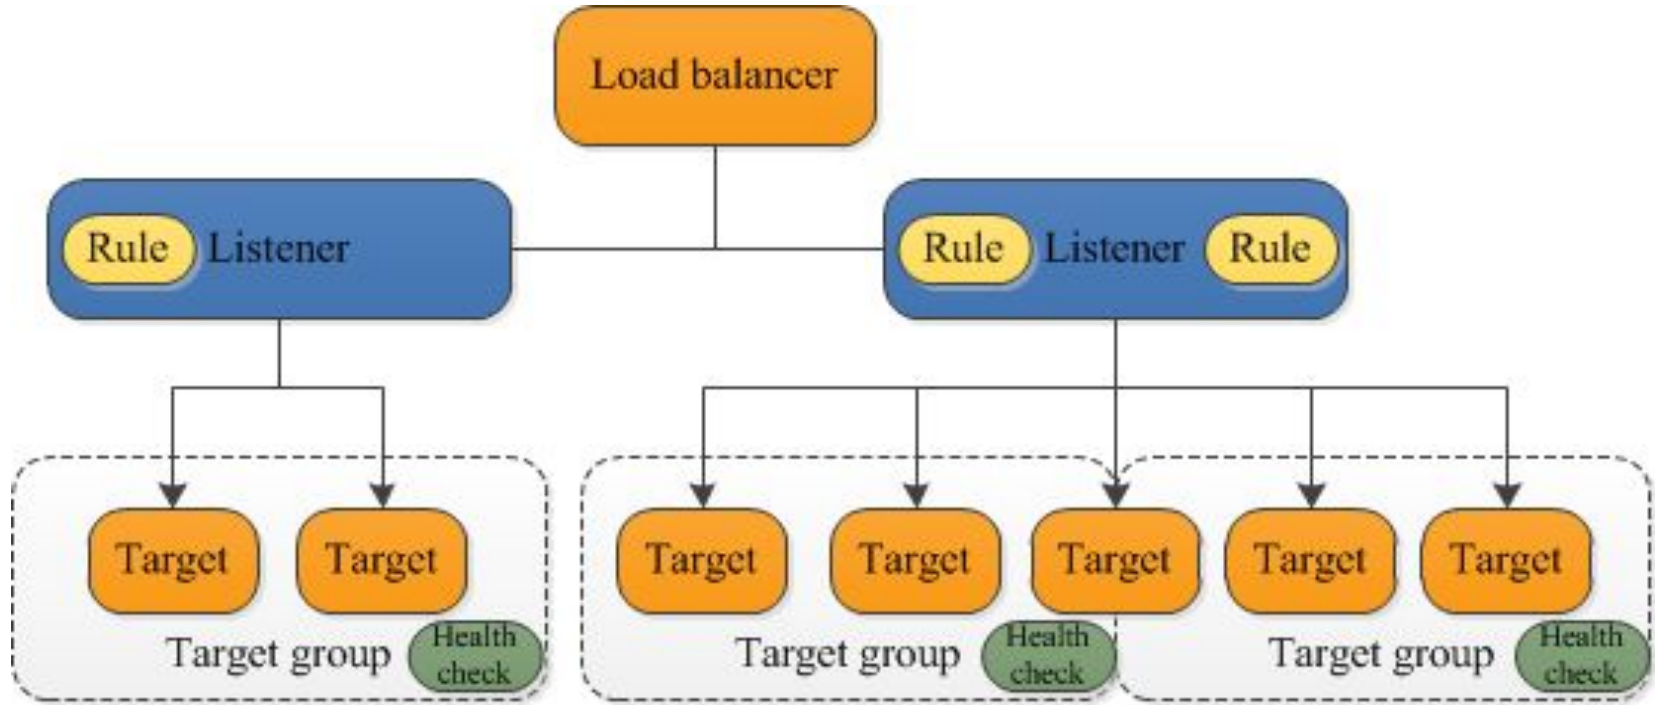
<https://aws.amazon.com/elasticloadbalancing/features/>

TCPI/IP and OSI model



<https://community.fs.com/blog/tcpip-vs-osi-whats-the-difference-between-the-two-models.html>

Application Load Balancer



Network Load Balancer

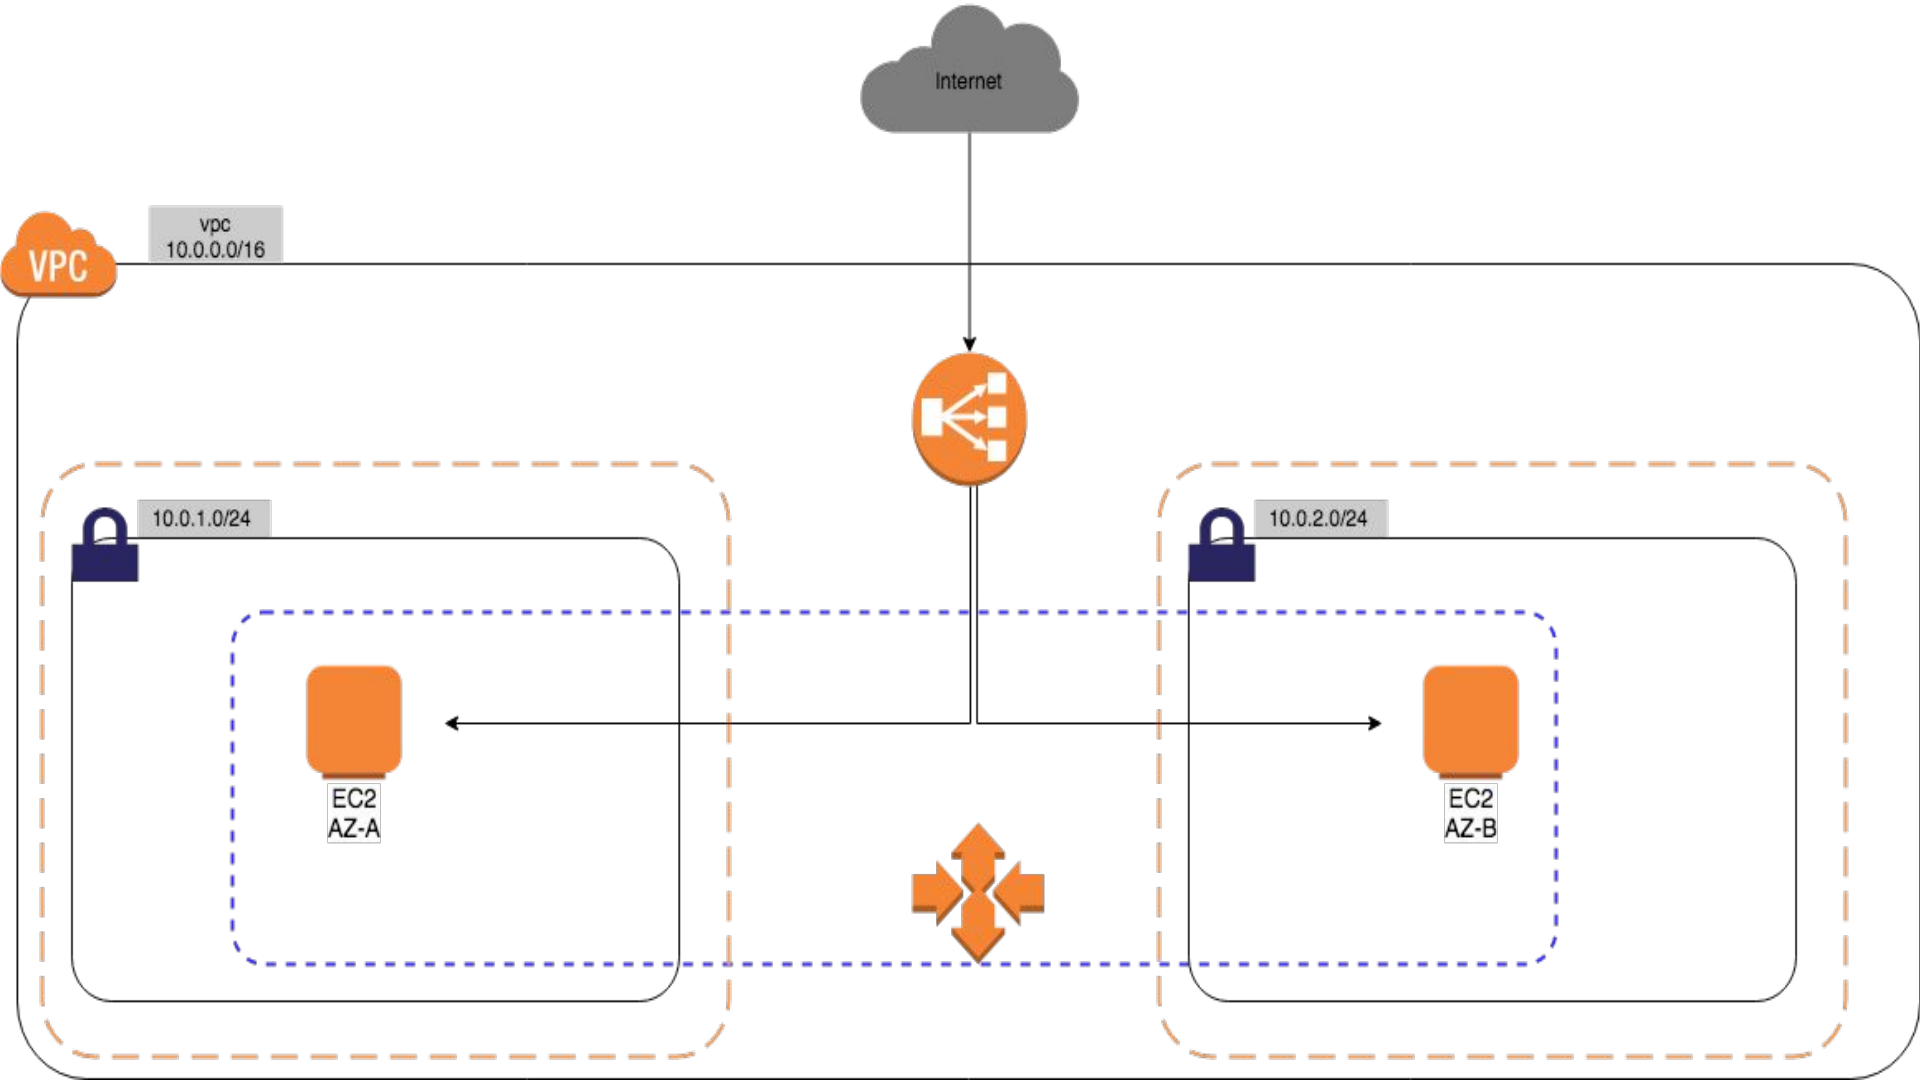
<https://docs.aws.amazon.com/elasticloadbalancing/latest/network/introduction.html>

Configure AWS Load Balancer

1. Create a Target Group
 - a. Protocol
 - i. TCP
 - ii. HTTP
 - iii. HTTPS
 - b. Target
 - i. Instance
 - ii. IP
2. Register instances(CLB)
3. Create ALB or NLB
4. Select the VPC and AZs
5. Configure security groups
6. Select the Target Group created earlier
7. Register targets

Auto Scaling groups for availability

- Components are:
 - EC2 instances are organized into groups for the purpose of scaling and management.
 - EC2 auto scaling group uses a launch configuration as a template for the instances.
 - Scaling plan: to tell Auto Scaling group when and how to scale.



Configure AWS Auto Scaling

From the console:

1. Create an Auto Scaling Group
2. Create launch configuration: the template used to launch auto scaling instances
 - a. Select a pre-configured AMI
 - b. Select type of instance
 - c. Attach IAM role (Best Practice)
 - d. Add storage if needed
3. Configure the auto scaling group details
 - a. Define the group size
 - b. Select the VPC
 - c. Associate Subnets
 - d. Use auto scaling group initial size or use policy
 - e. Configure Notification (optional)

Update Launch Configuration

To update launch configuration :

- 1) Copy the launch configuration and choose new name
 - a) Update the user data
 - b) Make sure all config is ok
- 2) Replace the Autoscaling with the new Launch Configuration
- 3) Remove instances: for each instance in the LB
 - a) terminate the instance
 - b) wait for the new instance to come up.
 - c) Test the new instance functionality
- 4) Remove the old launch configuration

CloudFormation

- Infrastructure as code: to build an automated delivery process for infrastructure, and updating it as needed.
- Express AWS resources in a text file in either a JSON format or YAML format.
- All template building blocks are optional except for “Resources”.
- Each template create a cloudformation stack

Cloudformation Template blocks

```
{
  "AWSTemplateFormatVersion" : "version date",
  "Description" : "JSON string",
  "Metadata" : {
    template metadata
  },
  "Parameters" : {
    set of parameters
  },
  "InstanceType" : {
    "Type" : "String",
    "Default" : "m1.small",
    "Description" : "EC2 instance type, e.g. m1.small, m1.large, etc."
  }
}
```

The latest template format version is 2010-09-09

String between 0 and 1024 bytes

Implementation details about specific resources

e.g. {"Description" : "Information about the instances"}

Create custom values as input when creating a stack

```
"Mappings" : {  
  set of mappings
```

Matches a key to a corresponding set of named values

e.g.

```
"RegionMap" : {  
  "us-east-1" : { "HVM64" : "ami-0ff8a91507f77f867" }  
},
```

```
"Conditions" : {  
  set of conditions
```

Used to match parameter values with resources or properties if the condition is true.

```
},
```

```
"Resources" : {  
  set of resources
```

Specifies the stack resources and their properties

e.g.

```
"MyEC2Instance" : {  
  "Type" : "AWS::EC2::Instance",  
  "Properties" : {  
    "ImageId" : "ami-0ff8a91507f77f867"  
  }  
},
```

```
"Outputs" : {  
  set of outputs
```

Declares output values that you can import into other stacks

e.g.

```
"InstanceID" : {  
  "Description": "The Instance ID",  
  "Value" : { "Ref" : "EC2Instance" }  
}  
}
```

Benefits of Cloudformation

- Automatically create and delete resources in order
- Update resources by updating the template
- One can monitor events that take place while cloudformation stack building is taking place from the console Events tab.
- Protection from deleting resources is available
- Specify deletion policies in template for specific resources
- If anything goes wrong while creating, deleting, and updating Cloudformation will rollback to the previous state automatically

Create VPC Environment

AWSTemplateFormatVersion: 2010-09-09

Description: "AWS CloudFormation Template to create a VPC with one public subnet and a private on"

Mappings:

CIDRConfig:

VPC:

CIDR: 10.10.0.0/16

Public:

CIDR: 10.10.0.0/24

Private:

CIDR: 10.10.128.0/17

Resources:

VPC:

Type: 'AWS::EC2::VPC'

Properties:

EnableDnsSupport: 'true'

EnableDnsHostnames: 'true'

CidrBlock: !FindInMap

- CIDRConfig
- VPC
- CIDR

Tags:

- Key: Name
- Value: !Sub '\${AWS::StackName}-vpc'

```
PublicSubnet:
  Type: 'AWS::EC2::Subnet'
  Properties:
    VpcId: !Ref VPC
    CidrBlock: !FindInMap
      - CIDRConfig
      - Public
      - CIDR
    Tags:
      - Key: Name
        Value: !Sub '${AWS::StackName}-public'
```

```
PrivateSubnet:
  Type: 'AWS::EC2::Subnet'
  Properties:
    VpcId: !Ref VPC
    CidrBlock: !FindInMap
      - CIDRConfig
      - Private
      - CIDR
    Tags:
      - Key: Name
        Value: !Sub '${AWS::StackName}-private'
```

InternetGateway:

Type: 'AWS::EC2::InternetGateway'

Properties:

Tags:

- Key: Name

Value: !Sub '\${AWS::StackName}-IGW'

GatewayToInternet:

Type: 'AWS::EC2::VPCGatewayAttachment'

Properties:

VpcId: !Ref VPC

InternetGatewayId: !Ref InternetGateway

NAT:

Type: 'AWS::EC2::NatGateway'

Properties:

AllocationId: !GetAtt

- EIP

- AllocationId

SubnetId: !Ref PublicSubnet

Tags:

- Key: Name

Value: !Sub '\${AWS::StackName}-NAT'

EIP:

Type: 'AWS::EC2::EIP'

Properties:

Domain: vpc

PublicRouteTable:

Type: 'AWS::EC2::RouteTable'

Properties:

VpcId: !Ref VPC

Tags:

- Key: Name

Value: !Sub '\${AWS::StackName}-publicRT'

PublicRoute:

Type: 'AWS::EC2::Route'

DependsOn: InternetGateway

Properties:

RouteTableId: !Ref PublicRouteTable

DestinationCidrBlock: 0.0.0.0/0

GatewayId: !Ref InternetGateway

PublicSubnetRouteTableAssociation:

Type: 'AWS::EC2::SubnetRouteTableAssociation'

Properties:

SubnetId: !Ref PublicSubnet

RouteTableId: !Ref PublicRouteTable

PrivateRouteTable:

Type: 'AWS::EC2::RouteTable'

Properties:

VpcId: !Ref VPC

Tags:

- Key: Name

Value: !Sub '\${AWS::StackName}-privateRT'

PrivateRoute:

Type: 'AWS::EC2::Route'

DependsOn: NAT

Properties:

RouteTableId: !Ref PrivateRouteTable

DestinationCidrBlock: 0.0.0.0/0

NatGatewayId: !Ref NAT

PrivateSubnetRouteTableAssociation:

Type: 'AWS::EC2::SubnetRouteTableAssociation'

Properties:

SubnetId: !Ref PrivateSubnet

RouteTableId: !Ref PrivateRouteTable

Outputs:

VPCId:

Description: VPCId of the newly created VPC

Value: !Ref VPC

Export:

Name: !Sub '\${AWS::StackName}-VPCID'

PublicSubnet:

Description: SubnetId of the public subnet

Value: !Ref PublicSubnet

Export:

Name: !Sub '\${AWS::StackName}-PublicSubID'

PrivateSubnet:

Description: SubnetId of the private subnet

Value: !Ref PrivateSubnet

Export:

Name: !Sub '\${AWS::StackName}-PrivateSubID'

Create EC2 Instance

AWSTemplateFormatVersion: 2010-09-09

Description: "AWS CloudFormation Template EC2InstanceWithSecurityGroup: Create an Amazon EC2"

Parameters:

 UserName:

 Type: String

 Default: e91

 SSHPublicKey:

 Type: String

 Default: ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAAA..

 InstanceType:

 Description: EC2 instance type

 Type: String

 Default: t2.micro

 AllowedValues:

- t2.micro
- t2.small
- t2.medium

 ConstraintDescription: must be a valid EC2 instance type.

 AMI:

 Description: AMI to create the EC2

 Type: String

 Default: ami-0dc7769bc7e889a35

 TagName:

 Description: Value of the Name tag for this EC2

 Type: String

 Default: e91

Resources:

SSHANDWEB:

Type: 'AWS::EC2::SecurityGroup'

Properties:

GroupName: SSH_LOCALWEB_INTERNET

GroupDescription: ssh access to my machine and 80 and 443 to the internet

VpcId: !ImportValue

'Fn::Sub': myStack1-VPCID

SecurityGroupIngress:

- IpProtocol: tcp

FromPort: '80'

ToPort: '80'

CidrIp: 0.0.0.0/0

- IpProtocol: tcp

FromPort: '443'

ToPort: '443'

CidrIp: 0.0.0.0/0

- IpProtocol: tcp

FromPort: '22'

ToPort: '22'

CidrIp: 140.247.0.0/16

nwFace1:

Type: 'AWS::EC2::NetworkInterface'

Properties:

SubnetId: !ImportValue

```
EC2Instance:
  Type: 'AWS::EC2::Instance'
  Properties:
    InstanceType: !Ref InstanceType
    NetworkInterfaces:
      - NetworkInterfaceId: !Ref nwFace1
        DeviceIndex: '0'
    ImageId: !Ref AMI
    Tags:
      -
        Key: Name
        Value: !Ref TagName
  UserData:
    Fn::Base64:
      Fn::Sub:
        - |
          #!/bin/bash -xe
          adduser ${USERNAME}
          echo ${USERNAME} 'ALL=(ALL) NOPASSWD:ALL'>>/etc/sudoers.d/${USERNAME}
          mkdir /home/${USERNAME}/.ssh
          echo ${SSHKEY} > /home/${USERNAME}/.ssh/authorized_keys
          chown -R ${USERNAME}.${USERNAME} /home/${USERNAME}/.ssh
          chmod 700 /home/${USERNAME}/.ssh
          chmod 600 /home/${USERNAME}/.ssh/authorized_keys
          hostnamectl set-hostname mystack1.harvard.edu --pretty
          hostnamectl set-hostname mystack1.harvard.edu --static
```



```
{
  "AWSTemplateFormatVersion" : "2010-09-09",

  "Description" : "This template creates a web server. **WARNING** This template creates an Amazon EC2 instance. You will be billed for the AWS resources used if you create a stack from this template.",

  "Parameters" : {

    "KeyName": {
      "Description" : "Name of an existing EC2 KeyPair to enable SSH access to the instances",
      "Type": "AWS::EC2::KeyPair::KeyName",
      "ConstraintDescription" : "must be the name of an existing EC2 KeyPair."
    },

    "InstanceType" : {
      "Description" : "WebServer EC2 instance type",
      "Type" : "String",
      "Default" : "t1.micro",
      "AllowedValues" : [ "t1.micro", "m3.medium", "m3.large", "c3.large", "c3.xlarge" ],
      "ConstraintDescription" : "must be a valid EC2 instance type."
    },

    "SSHLocation": {
      "Description": "The IP address range that can be used to SSH to the EC2 instances",
      "Type": "String",
      "MinLength": "9",
      "MaxLength": "18",
```

Monitoring and Logging

- What to monitor?
 - Availability - functional?
 - Security -
 - Performance - fast enough?
 -

Resources

Application Load Balancer:

<https://docs.aws.amazon.com/elasticloadbalancing/latest/application/introduction.html>

AS HealthCheck: <https://docs.aws.amazon.com/autoscaling/ec2/userguide/healthcheck.html>

ELB features: <https://aws.amazon.com/elasticloadbalancing/features/>